

Emotional Models: Supervised vs. Zero-Shot Classification

Miguel Castela¹ and Miguel Martins²

Univsersity of Coimbra, Portugal
uc2022212972@student.uc.pt
uc2022213951@student.uc.pt

1 Introduction

This project presents the implementation and evaluation of two distinct approaches to emotion classification in text using Transformer-based models applied to the *Emotion* dataset. The Emotion dataset, available on the HuggingFace Hub under the identifier `dair-ai/emotion`, consists of short English text samples primarily collected from social media platforms.

The dataset is pre-split into training, validation, and test sets. The main objective of this project is to compare the performance of two different methodologies for emotion classification. The main objective of this project is to compare the performance of two different methodologies for emotion classification. The first approach follows a supervised learning paradigm, where pretrained Transformer models are used as frozen feature extractors to generate sentence embeddings via models from the `sentence-transformers` library. The second approach adopts a zero-shot learning strategy based on Natural Language Inference (NLI) models, where emotion labels are expressed as natural language hypotheses and classification is performed without task-specific fine-tuning. By evaluating both supervised and zero-shot methods on the same dataset, this project aims to analyze their respective strengths and limitations, as well as characteristic error patterns.

Additionally, the impact of different pretrained models from the HuggingFace Hub is explored, assessing how variations in model size, architecture, and training objectives affect classification performance.

2 Description of the problem

The `dair-ai/emotion` dataset consists of text samples, primarily collected from Twitter, labeled with one of six emotions: *joy*, *sadness*, *anger*, *fear*, *love* and *surprise*. The training set containing 16,000 samples, the validation set 2,000 samples, and the test set 2,000 samples. The samples are distributed according to the following table:

Table 1. Class distribution of the `dair-ai/emotion` dataset.

Emotion	Train	Validation	Test	Total (%)
Joy	5,362	704	695	33.8%
Sadness	4,666	550	581	29.0%
Anger	2,159	275	275	13.5%
Fear	1,937	212	224	11.9%
Love	1,304	178	159	8.2%
Surprise	572	81	66	3.6%
Total	16,000	2,000	2,000	100.0%

By looking at table 1, it’s clear that the dataset is imbalanced, with the *joy* and *sadness* classes being significantly more represented, together accounting for over 60% of the data. In contrast, the *surprise* class is the least represented by a considerable margin, 3.6%. This imbalance could lead to biased model performance, favoring the majority classes, and makes it clear that “Accuracy” might not be the best metric to evaluate the models.

Another issue present is the mislabeling of some samples. After analysing the dataset, we found samples labeled with the wrong emotions. For example, the sentence “I didn’t feel humiliated” was labeled as *sadness*. The sentence explicitly denies the negative emotion, yet it was labeled with it, presumably due to the presence of the word “humiliated”. This kind of mislabeling will inevitably cap the performance of any model trained on this dataset, as it introduces noise and confusion during the learning process. Another possible cause is the fact that dataset relies on th hastags to label the samples, and their use in sarcastic or ironic contexts is a major source of error.

3 Approach of Supervised Classification

This approach follows a supervised learning paradigm in which pretrained Transformer models are used as fixed feature extractors, and a lightweight neural classifier is trained on top of the resulting sentence embeddings. The model leverages their pretrained semantic representations to reduce computational cost and mitigate overfitting, while still capturing rich contextual information from the input text.

The Emotion dataset is loaded using the HuggingFace `datasets` library and split into the training, validation, and test sets. To convert raw text into numerical representations suitable for classification, a pretrained sentence-level Transformer from the `sentence-transformers` library is employed. ned

3.1 Embedding Generation

Models are used to encode each sentence into a dense, fixed-length embedding vector. The embedding process is performed in batches to improve efficiency and introduces beneficial stochasticity into the optimization process. , and the resulting representations are stored as `NumPy` arrays for the training, validation, and test sets. The Transformer remains frozen during this stage, meaning that no gradient updates are applied to its parameters.

These embeddings serve as high-level semantic features that capture syntactic and contextual information from the input text. The corresponding emotion labels are retained as integer class indices, resulting in pairs of the form (*embedding*, *label*), which are later wrapped into `PyTorch DataLoader` objects for mini-batch training and evaluation.

On top of the generated embeddings, we trained and evaluated two light classification heads, Multi-Layer Perceptron and Logistic Regression.

3.2 Classifier Architecture (ANN)

Equivalent to a small Multi-Layer Perceptron, its architecture is intentionally lightweight, as the main representational power is delegated to the pretrained Transformer embeddings. The architecture, defined in our `ANNClassifier` module, consists of:

- **Input Layer:** Matches the dimension of the embeddings.
- **Hidden Layer:** A fully connected layer with 256 neurons, followed by a Rectified Linear Unit (ReLU) activation function, introducing non-linearity into the model.
- **Regularization:** A Dropout layer, implemented in the training, with a probability of $p = 0.1$ to prevent overfitting.
- **Output Layer:** A fully connected layer with six neurons, one for each emotion class, producing unnormalized logits.

Class predictions are obtained by applying the `argmax` operation over the output logits, selecting the emotion class with the highest predicted score.

3.3 Classifier Architecture (Logistic Regression)

As an alternative to the neural classifier, a multinomial Logistic Regression (LogReg) model is employed as a lightweight linear classification head on top of the frozen Transformer embeddings. This approach serves as a strong and interpretable baseline, allowing us to assess how well linearly separable the embedding space is for the emotion classification task.

The Logistic Regression classifier is implemented using the `scikit-learn` library and trained directly on the precomputed sentence embeddings, without any additional feature transformation or hidden layers. Its architecture can be summarized as follows:

- **Input Layer:** A fixed-size embedding vector whose dimensionality matches the output of the pretrained Transformer model.
- **Linear Decision Function:** A single linear transformation that maps the input embeddings to six output scores, one for each emotion class.
- **Multinomial Softmax:** The class probabilities are obtained via the multinomial logistic `softmax` formulation.

4 Approach of Zero-Shot Classification

The zero-shot classification approach leverages pre-trained Natural Language Inference (NLI) models to classify text without specific task-related training data. NLI models are trained to determine the relationship between a premise and a hypothesis. The classification setup requires each candidate label to be converted into a natural language hypothesis. For this purpose, the template "This text expresses ." is used (NLI.py).

4.1 Architecture and Inference Procedure

1. **Batch Inference:** The test texts are processed in batches for computational efficiency.
2. **Prediction:** For each text, the model compares the text (premise) against a hypothesis for every candidate label.
3. **Label Selection:** The label whose hypothesis yields the highest entailment score is selected as the predicted class for the input text.

5 Experimental Setup

In this section, we detail the specific models employed in both the supervised and zero-shot classification approaches, along with their configurations and relevant hyperparameters.

5.1 ANN

Training Procedure

The classifier is trained using the **Cross-Entropy** loss function, which is standard for multi-class classification tasks.

To prevent overfitting and ensure robust model selection, **early stopping** is employed based on the macro-averaged F1 score computed on the validation set.

Training is halted if no improvement in validation F1 is observed for a fixed number of consecutive epochs, and the best-performing model state is restored before final evaluation.

The ANN utilized in supervised training was configured with the hyperparameters listed in Table 2. The `all-MiniLM-L6-v2`, `paraphrase-mpnet-base-v2` and `all-mpnet-base-v2` models from the `sentence-transformers` library were employed to generate sentence embeddings. Their configurations, including embedding dimensions and batch sizes, are summarized in Table 3. Experiments were conducted using five different seeds to account for variability in training outcomes.

Table 2. Hyperparameters for supervised training

Hyperparameter	Value
Batch Size	64
Learning Rate	1×10^{-3}
Epochs	25 with Early Stopping
Loss Function	Cross-Entropy Loss
Optimizer	Adam
Seeds	[2025, 2029]

Table 3. Supervised Models and their configurations.

Model	Embedding Dimension
all-MiniLM-L6-v2	384
paraphrase-mpnet-base-v2	768
all-mpnet-base-v2	768

Evaluation

After training, the selected model is evaluated on the held-out test set. Performance is measured using accuracy, macro-averaged F1 score, and confusion matrices. To account for randomness introduced by weight initialization and data shuffling, experiments are repeated across multiple random seeds, and results are reported as mean values with corresponding standard deviations. This evaluation protocol enables a fair and reliable assessment of the supervised approach based on frozen Transformer embeddings.

6 Logistic Regression

Training Procedure

The classifier is trained using the `lbfgs` optimizer with a maximum of 200 it-

erations and the **multinomial** setting, which jointly optimizes all class decision boundaries.

Unlike the ANN-based classifier, the Logistic Regression head does not require iterative epoch-based training or early stopping. The optimization problem is convex, and training converges deterministically to a global optimum given the same input data. As a result, the model is trained once per embedding configuration using the full training set.

Although experiments are repeated across multiple random seeds for consistency with the ANN evaluation protocol, reported metrics for this classifier exhibit zero standard deviation across runs. This behavior is expected, as Logistic Regression is a deterministic linear model and when trained on identical embeddings and labels, it yields the same learned parameters and predictions regardless of the random seed.

Evaluation

After training, the Logistic Regression classifier is evaluated on the held-out test set using accuracy, macro-averaged F1 score, and confusion matrices. Predictions are obtained by selecting the class with the highest predicted probability. The deterministic nature of the model allows for a direct and stable comparison with the ANN classifier, highlighting the performance gains introduced by non-linear classification layers on top of the same frozen embeddings.

7 NLI

As described in Section 3.1, the zero-shot classification approach utilizes pre-trained NLI models to perform emotion classification without task-specific training data. The models evaluated in this approach are summarized in Table 4. All models were sourced from the HuggingFace Hub and are tested with a batch size of 32 on a GPU device to optimize inference speed.

Table 4. NLI Models used for Zero-Shot Classification

Model Name	Size (parameters)
MoritzLaurer/xtremedistil-l6-h256-zeroshot-v1.1-all-33	$\approx$ 13M
joeddav/xlm-roberta-large-xnli	$\approx$ 600M
MoritzLaurer/bge-m3-zeroshot-v2.0	$\approx$ 600M
MoritzLaurer/deberta-v3-xsmall-zeroshot-v1.1-all-33	$\approx$ 71M
MoritzLaurer/ModernBERT-large-zeroshot-v2.0	$\approx$ 400M
facebook/bart-large-mnli	$\approx$ 400M

Evaluation Metrics

Unlike the supervised pipeline, this method was applied directly to the test set

without any parameter updates. The performance of the zero-shot classifier is evaluated on the test set using standard classification metrics. This ensures a fair comparison with the supervised approach. The primary metrics used for evaluation are accuracy, macro-averaged F1 score and confusion matrices, which provide insights into the model’s overall performance and its ability to handle class imbalances. This evaluation framework allows for a comprehensive comparison between the zero-shot and supervised classification approaches.

8 Results and Discussion

8.1 Supervised Models

The following table (Table 5) summarizes the performance across the different embedding models and classifier heads used in the supervised approach. Results are reported as mean accuracy and mean macro F1 score over five different random seeds, along with the average number of epochs taken for training before early stopping.

Model	Head	Accuracy	Macro F1	Avg. Epochs
all-MiniLM-L6-v2	ANN	0.7315 ± 0.0061	0.6556 ± 0.0057	20.0
all-MiniLM-L6-v2	LogReg	0.6870 ± 0.0000	0.5928 ± 0.0000	N/A
paraphrase-mpnet-base-v2	ANN	0.7089 ± 0.0037	0.6249 ± 0.0091	12.4
paraphrase-mpnet-base-v2	LogReg	0.6950 ± 0.0000	0.6083 ± 0.0000	N/A
all-mpnet-base-v2	ANN	0.7014 ± 0.0019	0.6177 ± 0.0059	13.0
all-mpnet-base-v2	LogReg	0.6760 ± 0.0000	0.5629 ± 0.0000	N/A

Table 5. Performance Comparison of Embedding Models and Classifiers

Classifier Comparison: The ANN classifier consistently outperforms the Logistic Regression head across all embedding models. Even though this gap is smaller in `paraphrase-mpnet-base-v2` model, it made a noticeable difference in the other two. This indicates that the non-linear decision boundaries learned by the ANN provide a significant advantage over the linear separability assumed by Logistic Regression, and suggests that the embedding spaces are not linearly separable for the emotion classification task.

It is also noticeable that apart from the `all-MiniLM-L6-v2` model which on average used all 25 epochs, early stopping occurred relatively quickly on the rest, within the first 15 epochs. This suggests that it might’ve been beneficial to train the first model for longer, as the scores were still rising

Model Comparison: Among the embedding models evaluated, the highest performance scores were achieved by `all-MiniLM-L6-v2` in both accuracy and

macro F1 score, regardless of the classifier head used. However, the performance differences between the embedding models are not substantial, indicating that all three models provide reasonably effective representations for the emotion classification task. The size of the models did also not show much effect on the performance, as the smallest model (with dimension 384) actually performed the best in comparison to the larger ones (with dimension 768).

8.2 NLI Models

The following table (Table 6) summarizes the performance of the different NLI models used in the zero-shot classification approach. Results are reported in terms of accuracy and macro F1 score on the test set.

Table 6. Zero-Shot Classification Results

Model Name	ACC	Macro F1
MoritzLaurer/xtremedistil-l6-h256-zeroshot-v1.1-all-33	0.6120	0.5148
joeddav/xlm-roberta-large-xnli	0.3505	0.3222
MoritzLaurer/bge-m3-zeroshot-v2.0	0.6960	0.6232
MoritzLaurer/deberta-v3-xsmall-zeroshot-v1.1-all-33	0.7045	0.6306
MoritzLaurer/ModernBERT-large-zeroshot-v2.0	0.7090	0.6304
facebook/bart-large-mnli	0.4660	0.4333

Model Comparison: The results indicate significant variability in performance across the different NLI models. The model that achieved the highest macro F1 score of 0.6304 was `MoritzLaurer/ModernBERT-large-zeroshot-v2.0`, closely followed by `MoritzLaurer/deberta-v3-xsmall-zeroshot-v1.1-all-33` with a macro F1 of 0.6306. The worst performing model achieved a macro F1 score of only 0.3222 and was `joeddav/xlm-roberta-large-xnli`.

Interestingly, model size did not correlate strongly with model performance. For example, `MoritzLaurer/deberta-v3-xsmall-zeroshot-v1.1-all-33`, which is relatively small, with ≈ 71 M parameters, outperformed the largest models such as `joeddav/xlm-roberta-large-xnli` with ≈ 600 M parameters. Models with similar sizes also exhibited varying performance, like `facebook/bart-large-mnli` and `MoritzLaurer/ModernBERT-large-zeroshot-v2.0`, both around ≈ 400 M parameters.

This discrepancy is due to a different reason: what each model was taught during pretraining. Taking `joeddav/xlm-roberta-large-xnli` and `deberta` as examples. On the one hand, `xlm-roberta` was trained in a multilingual setting, which makes it less specialized in English tasks, and on standard NLI datasets, focusing on logic reasoning. On the other hand, `deberta` was trained in a monolingual setting, specifically for zero-shot classification tasks, and its training included emotion detection, sentiment, and topic classification, which is far more relevant

to our task. This proved to be a major factor when it came to performance and that, given a different task, the results could be completely different. This difference could also be attributed to the architecture of each model, as some may already have an outdated architecture that doesn't perform as well as more recent ones. For example, `facebook/bart-large-mnli` is based on the BART architecture, which, while powerful, has been surpassed by newer architectures like DeBERTa and ModernBERT in various NLP tasks.

8.3 Supervised vs Zero-Shot Classification

8.4 NLI Models




Fig. 1. A figure caption is always placed below the illustration. Please note that short captions are centered, while long ones are justified by the macro package automatically.

Theorem 1. *This is a sample theorem. The run-in heading is set in bold, while the following text appears in italics. Definitions, lemmas, propositions, and corollaries are styled the same way.*

Proof. Proofs, examples, and remarks have the initial word in italics, while the following text appears in normal font.

For citations of references, we prefer the use of square brackets and consecutive numbers. Citations using labels or the author/year convention are also acceptable. The following bibliography provides a sample reference list with entries for journal articles [1], an LNCS chapter [2], a book [3], proceedings without editors [4], and a homepage [5]. Multiple citations are grouped [1–3], [1, 3–5].

Acknowledgments. A bold run-in heading in small font size at the end of the paper is used for general acknowledgments, for example: This study was funded by X (grant number Y).

Disclosure of Interests. It is now necessary to declare any competing interests or to specifically state that the authors have no competing interests. Please place the statement with a bold run-in heading in small font size beneath the (optional) acknowledgments¹, for example: The authors have no competing interests to declare that are relevant to the content of this article. Or: Author A has received research grants from Company W. Author B has received a speaker honorarium from Company X and owns stock in Company Y. Author C is a member of committee Z.

References

1. Author, F.: Article title. Journal **2**(5), 99–110 (2016)
2. Author, F., Author, S.: Title of a proceedings paper. In: Editor, F., Editor, S. (eds.) CONFERENCE 2016, LNCS, vol. 9999, pp. 1–13. Springer, Heidelberg (2016). <https://doi.org/10.1007/1234567890>
3. Author, F., Author, S., Author, T.: Book title. 2nd edn. Publisher, Location (1999)
4. Author, A.-B.: Contribution title. In: 9th International Proceedings on Proceedings, pp. 1–2. Publisher, Location (2010)
5. LNCS Homepage, <http://www.springer.com/lncs>, last accessed 2023/10/25

¹ If EquinOCS, our proceedings submission system, is used, then the disclaimer can be provided directly in the system.